

simplelinearregression

August 25, 2023

1 Simple Linear Regression

- Linear Regression is used for predicting the relationship between the independent and dependent variable. It aims to predict the good fit line that describes the relationship by minimizing the mean_squared_error of the predicted value and the actual value.
- In simple linear regression, we only use one independent variable or feature to predict the target variable.
- Slope-intercept formulae,
 $Y = B_0 + B_1X$ where, y = target variable B_0 = intercept B_1 = slope/coefficient X = Predictive variable
- Linear regression aims to find the slope value and intercept value to predict the target feature.
- Have used the Tv-Marketing data set from kaggle.com; where, Independent variable(X)- TV and target variable(Y)- Sales.

```
[ ]: # Importing necessary libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[ ]: # Onboarding data

from google.colab import files
rawdata=files.upload()
```

<IPython.core.display.HTML object>

Saving tvmarketing.csv to tvmarketing.csv

```
[ ]: # DataFrame

df=pd.read_csv('tvmarketing.csv')
df
```

```
[ ]:      TV  Sales
0    230.1   22.1
1     44.5   10.4
2     17.2    9.3
3    151.5   18.5
4    180.8   12.9
..     ...   ...
195   38.2    7.6
196   94.2    9.7
197  177.0   12.8
198  283.6   25.5
199  232.1   13.4

[200 rows x 2 columns]
```

```
[ ]: # Shallow copy

df_copy=df.copy()
```

2 Exploratory Data Analysis

```
[ ]: df.head()
```

```
[ ]:      TV  Sales
0    230.1   22.1
1     44.5   10.4
2     17.2    9.3
3    151.5   18.5
4    180.8   12.9
```

```
[ ]: df.shape
```

```
[ ]: (200, 2)
```

```
[ ]: df.columns
```

```
[ ]: Index(['TV', 'Sales'], dtype='object')
```

```
[ ]: # Technical Report

df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype

```

```

-----
0   TV      200 non-null   float64
1   Sales   200 non-null   float64
dtypes: float64(2)
memory usage: 3.2 KB

```

```
[ ]: # Statistical Report
```


```
df.describe()
```

```
[ ]:
      TV      Sales
count 200.000000 200.000000
mean  147.042500  14.022500
std    85.854236   5.217457
min     0.700000   1.600000
25%    74.375000  10.375000
50%   149.750000  12.900000
75%   218.825000  17.400000
max   296.400000  27.000000


```

2.1 Univariate Analysis


```
[ ]: sns.boxplot(df['TV'],color='pink')
plt.show()
```



```
[ ]: sns.scatterplot(df,x='TV',y='Sales',size=5,color='green')  
plt.show()
```



```
[ ]: sns.heatmap(df.corr())  
plt.show()
```



```
[ ]: # Null values
df.isna().sum()/len(df)*100
```

```
[ ]: TV      0.0
Sales    0.0
dtype: float64
```

```
[ ]: # Duplicate values
df.duplicated().sum()
```

```
[ ]: 0
```

3 Standardization

```
[ ]: X=df.iloc[:,[0]]
Y=df.iloc[:,[-1]]
```

```
[ ]: from sklearn.preprocessing import StandardScaler
```

```
sc=StandardScaler()  
X_sc=sc.fit_transform(X)
```

4 Train Test Split

```
[ ]: from sklearn.model_selection import train_test_split  
  
x_train,x_test,y_train,y_test=train_test_split(X_sc,Y,test_size=0.  
↳3,random_state=4)
```

```
[ ]: from sklearn.linear_model import LinearRegression  
  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

```
[ ]: LinearRegression()
```

```
[ ]: y_predict=lr.predict(x_test)
```

```
[ ]: y_predict
```

```
[ ]: array([[17.16400307],  
           [13.33072137],  
           [17.4340204 ],  
           [20.1631241 ],  
           [ 8.95740501],  
           [18.46587233],  
           [13.28732572],  
           [15.01832966],  
           [11.06932625],  
           [16.0260729 ],  
           [19.41093298],  
           [18.00298548],  
           [10.45214378],  
           [12.74729107],  
           [19.04930263],  
           [20.60190226],  
           [19.61826771],  
           [14.7434906 ],  
           [ 8.88507894],  
           [17.10614222],  
           [18.37425931],  
           [20.32224146],  
           [18.5381984 ],  
           [11.42613486],  
           [18.37908105],
```

```

[17.87762029],
[15.28352525],
[ 9.58423094],
[16.65289884],
[ 8.71631811],
[10.49553943],
[10.14837429],
[16.78308577],
[15.25459483],
[15.07136878],
[ 7.74714877],
[17.5352769 ],
[10.10980039],
[17.34722912],
[12.10599992],
[13.70199519],
[ 8.59095292],
[14.3625733 ],
[ 7.37587494],
[ 7.7133966 ],
[12.10599992],
[17.76672032],
[13.52359089],
[ 7.95930524],
[17.26525957],
[17.69921599],
[15.86695555],
[10.08086996],
[13.91415166],
[ 7.64107053],
[11.85526954],
[16.69629449],
[10.21587862],
[ 7.00942286],
[ 9.81567437]])

```

```
[ ]: y_test
```

```

[ ]: Sales
11    17.4
99    17.2
128   24.7
175   27.0
1     10.4
111   21.8
90    11.2
177   11.7

```

88	12.9
187	17.3
61	24.2
199	13.4
191	9.9
123	15.2
184	17.6
188	15.9
33	17.4
171	14.5
138	9.6
84	21.7
81	12.3
102	14.8
147	25.4
34	9.5
47	23.2
124	19.7
112	14.1
6	11.8
14	19.0
190	10.8
80	11.8
18	11.3
167	12.2
45	14.9
153	19.0
119	6.6
100	11.7
83	13.6
181	12.2
71	12.4
26	15.0
134	10.8
180	10.5
158	7.3
189	6.7
89	16.7
48	14.8
116	12.2
12	9.2
69	22.3
110	13.4
154	15.6
16	12.5
19	14.6
2	9.3


```
143    10.4
185    22.6
29     10.5
155     3.2
24     9.7
```

```
[ ]: # Intercept value
```

```
lr.intercept_
```

```
[ ]: array([13.90173569])
```

```
[ ]: # Slope value
```

```
lr.coef_
```

```
[ ]: array([[4.1293042]])
```

5 Model Evaluation

```
[ ]: from sklearn.metrics import mean_squared_error,mean_absolute_error,r2_score
```

```
MAE=mean_absolute_error(y_test,y_predict)
MSE=mean_squared_error(y_test,y_predict)
RMSE=np.sqrt(MSE)
print("MAE",MAE)
print("MSE",MSE)
print("RMSE",RMSE)
```

```
MAE 2.7404123217291905
MSE 11.404170638824256
RMSE 3.3770061650557075
```

```
[ ]: R2=r2_score(y_test,y_predict)
print("R2",R2)
```

```
R2 0.5524131166103536
```

The R2 score reveals that the model accuracy is 55% (i.e) 55% of the data fits in the good fit line/regression line.

6 Data visualization

```
[ ]: plt.scatter(x_train,y_train,color='blue')
plt.plot(x_train,lr.predict(x_train),color='red')
plt.title('Sales Prediction')
plt.xlabel('TV')
```

```
plt.ylabel('Sales')  
plt.grid()  
plt.show()
```

